

Quiz7-4

```
library(ggplot2)
set.seed(123)
n <- 1000

year_of_construction <- sample(1700:2020, n, replace = TRUE)
location_zone <- sample(c("Central", "Suburban", "Outskirts"), n, replace = TRUE, prob = c(0.3, 0.4, 0.3))
building_type <- sample(c("Residential", "Commercial", "Mixed-use"), n, replace = TRUE)

number_of_floors <- round(runif(n, min = 1, max = 100) *
                           (year_of_construction / 2020)^2 *
                           (ifelse(location_zone == "Central", 1.5, 1)) *
                           (ifelse(building_type == "Commercial", 1.2, 1))
                           )

buildings_df <- data.frame(year_of_construction, location_zone, building_type, number_of_floors)


head(buildings_df)
```

	year_of_construction	location_zone	building_type	number_of_floors
1	1878	Suburban	Mixed-use	30
2	1713	Central	Commercial	34
3	1894	Suburban	Residential	23
4	2005	Central	Mixed-use	116
5	1817	Central	Commercial	27
6	1998	Suburban	Residential	59

```
ggplot(buildings_df, aes(x = year_of_construction, y = number_of_floors)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "lm", se = FALSE, color = "blue") +
```

```
labs(title = "Relationship Between Year of Construction and Number of Floors",
     x = "Year of Construction",
     y = "Number of Floors") +
theme_minimal()
```

`geom_smooth()` using formula = 'y ~ x'



```
if (!require("rstanarm")) install.packages("rstanarm")
```

Loading required package: rstanarm

Loading required package: Rcpp

This is rstanarm version 2.32.1

- See <https://mc-stan.org/rstanarm/articles/priors> for changes to default priors!
- Default priors may change, so it's safest to specify priors, even if equivalent to the default
- For execution on a local, multicore CPU with excess RAM we recommend calling

```
options(mc.cores = parallel::detectCores())
```

```
library(rstanarm)

model <- stan_glm(number_of_floors ~ year_of_construction,
  data = buildings_df,
  family = gaussian, # Because we are predicting a continuous outcome
  prior = normal(0, 2.5), # Assuming a normal prior with mean 0 and SD 2.5
  seed = 123 # Set a seed for reproducibility
)
```

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 2.7e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.27 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)

Chain 1: Iteration: 200 / 2000 [10%] (Warmup)

Chain 1: Iteration: 400 / 2000 [20%] (Warmup)

Chain 1: Iteration: 600 / 2000 [30%] (Warmup)

Chain 1: Iteration: 800 / 2000 [40%] (Warmup)

Chain 1: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 1: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 1: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 1: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 1: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 1: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.036 seconds (Warm-up)

Chain 1: 0.095 seconds (Sampling)

Chain 1: 0.131 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 9e-06 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.09 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 2000 [0%] (Warmup)

Chain 2: Iteration: 200 / 2000 [10%] (Warmup)

Chain 2: Iteration: 400 / 2000 [20%] (Warmup)

Chain 2: Iteration: 600 / 2000 [30%] (Warmup)

Chain 2: Iteration: 800 / 2000 [40%] (Warmup)

Chain 2: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 2: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 2: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 2: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 2: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 2: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 2:

Chain 2: Elapsed Time: 0.032 seconds (Warm-up)

Chain 2: 0.096 seconds (Sampling)

Chain 2: 0.128 seconds (Total)

Chain 2:

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).

Chain 3:

Chain 3: Gradient evaluation took 1e-05 seconds

Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.1 seconds.

Chain 3: Adjust your expectations accordingly!

Chain 3:

Chain 3:

Chain 3: Iteration: 1 / 2000 [0%] (Warmup)

Chain 3: Iteration: 200 / 2000 [10%] (Warmup)

Chain 3: Iteration: 400 / 2000 [20%] (Warmup)

Chain 3: Iteration: 600 / 2000 [30%] (Warmup)

Chain 3: Iteration: 800 / 2000 [40%] (Warmup)

Chain 3: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 3: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 3: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 3: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 3: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 3: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 3:

Chain 3: Elapsed Time: 0.07 seconds (Warm-up)

Chain 3: 0.093 seconds (Sampling)

Chain 3: 0.163 seconds (Total)

Chain 3:

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).

Chain 4:

Chain 4: Gradient evaluation took 1e-05 seconds

Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.1 seconds.

Chain 4: Adjust your expectations accordingly!

Chain 4:

Chain 4:

Chain 4: Iteration: 1 / 2000 [0%] (Warmup)

Chain 4: Iteration: 200 / 2000 [10%] (Warmup)

Chain 4: Iteration: 400 / 2000 [20%] (Warmup)

Chain 4: Iteration: 600 / 2000 [30%] (Warmup)

Chain 4: Iteration: 800 / 2000 [40%] (Warmup)

Chain 4: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 4: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 4: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 4: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 4: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 4: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 4:

Chain 4: Elapsed Time: 0.039 seconds (Warm-up)

Chain 4: 0.09 seconds (Sampling)

Chain 4: 0.129 seconds (Total)

Chain 4:

```
# View model summary
print(summary(model))
```

Model Info:

```
function: stan_glm
family: gaussian [identity]
formula: number_of_floors ~ year_of_construction
algorithm: sampling
sample: 4000 (posterior sample size)
priors: see help('prior_summary')
observations: 1000
predictors: 2
```

Estimates:

	mean	sd	10%	50%	90%
(Intercept)	-53.8	20.1	-79.5	-54.0	-28.1
year_of_construction	0.1	0.0	0.0	0.1	0.1
sigma	31.5	0.7	30.6	31.5	32.4

Fit Diagnostics:

	mean	sd	10%	50%	90%
mean_PPD	51.4	1.4	49.6	51.4	53.2

The mean_ppd is the sample average posterior predictive distribution of the outcome variable

MCMC diagnostics

	mcse	Rhat	n_eff
(Intercept)	0.3	1.0	4285
year_of_construction	0.0	1.0	4283
sigma	0.0	1.0	3351
mean_PPD	0.0	1.0	3043
log-posterior	0.0	1.0	2120

For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective

```
# Visualize model results
plot(model)
```

